

Deep Learning


Convolutional Neural Networks

What are Hidden Layers Doing?

Learning Feature Detectors

Sub Question 1: What sort of features can the neurons in a first hidden layer learn in a fully connected network?

Sub Question 2: Could a Hidden Layer learn to recognize a specific black and white face in a specific position in an image?



Problem 1: Computational Explosion

- The number of weights to be learned in a fully connected layer between input units L1 and hidden units L2 is:

$$(size(L1) + 1) \times size(L2)$$

- How many weights (or parameters) are required for a 256 x 256 three colour image that is to be fully connected to a hidden layer with 128 neurons?

25 Million!!!!

Problem 2: Local vs. Global Detection



Interesting things in images are usually **translation invariant**

In early layers we can focus on local features and leave global features until later

Edge Detection – Example of a Local Feature detector that can be applied to a whole image

An edge is a steep gradient in intensities


-1	0	+1
-1	0	+1
-1	0	+1

Gx

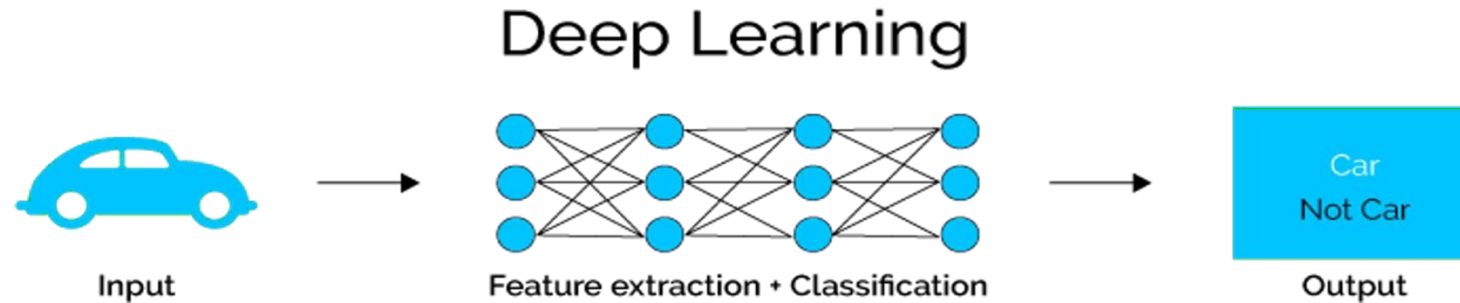
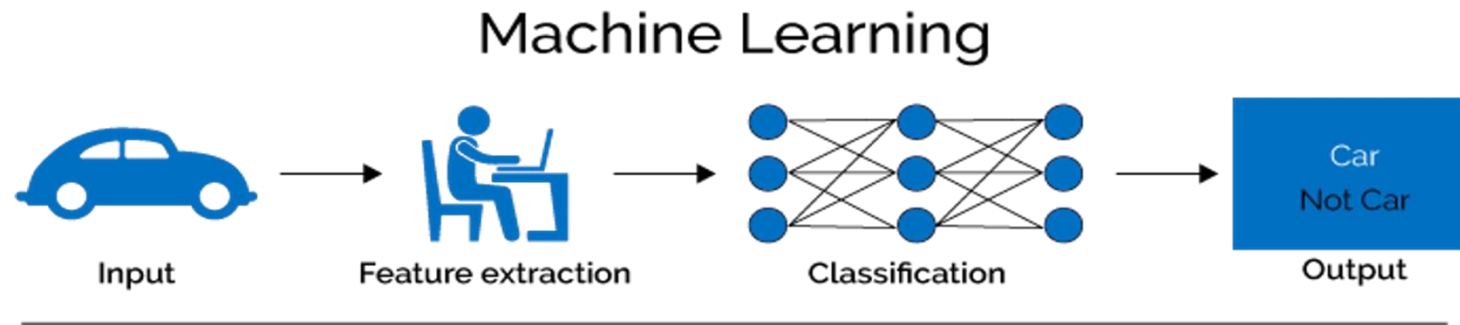
+1	+1	+1
0	0	0
-1	-1	-1

Gy

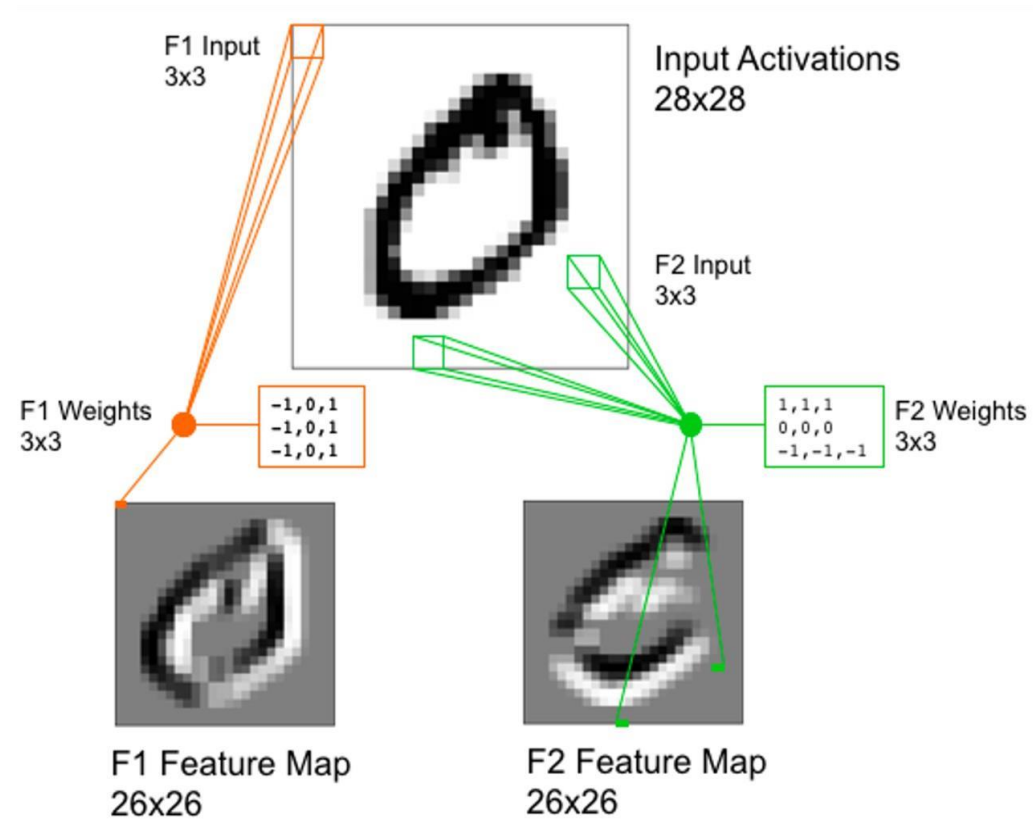
*Weight matrix for
edge detection*




CNNs - Let's Automate This




Edge Detectors, Filters, Features and Maps



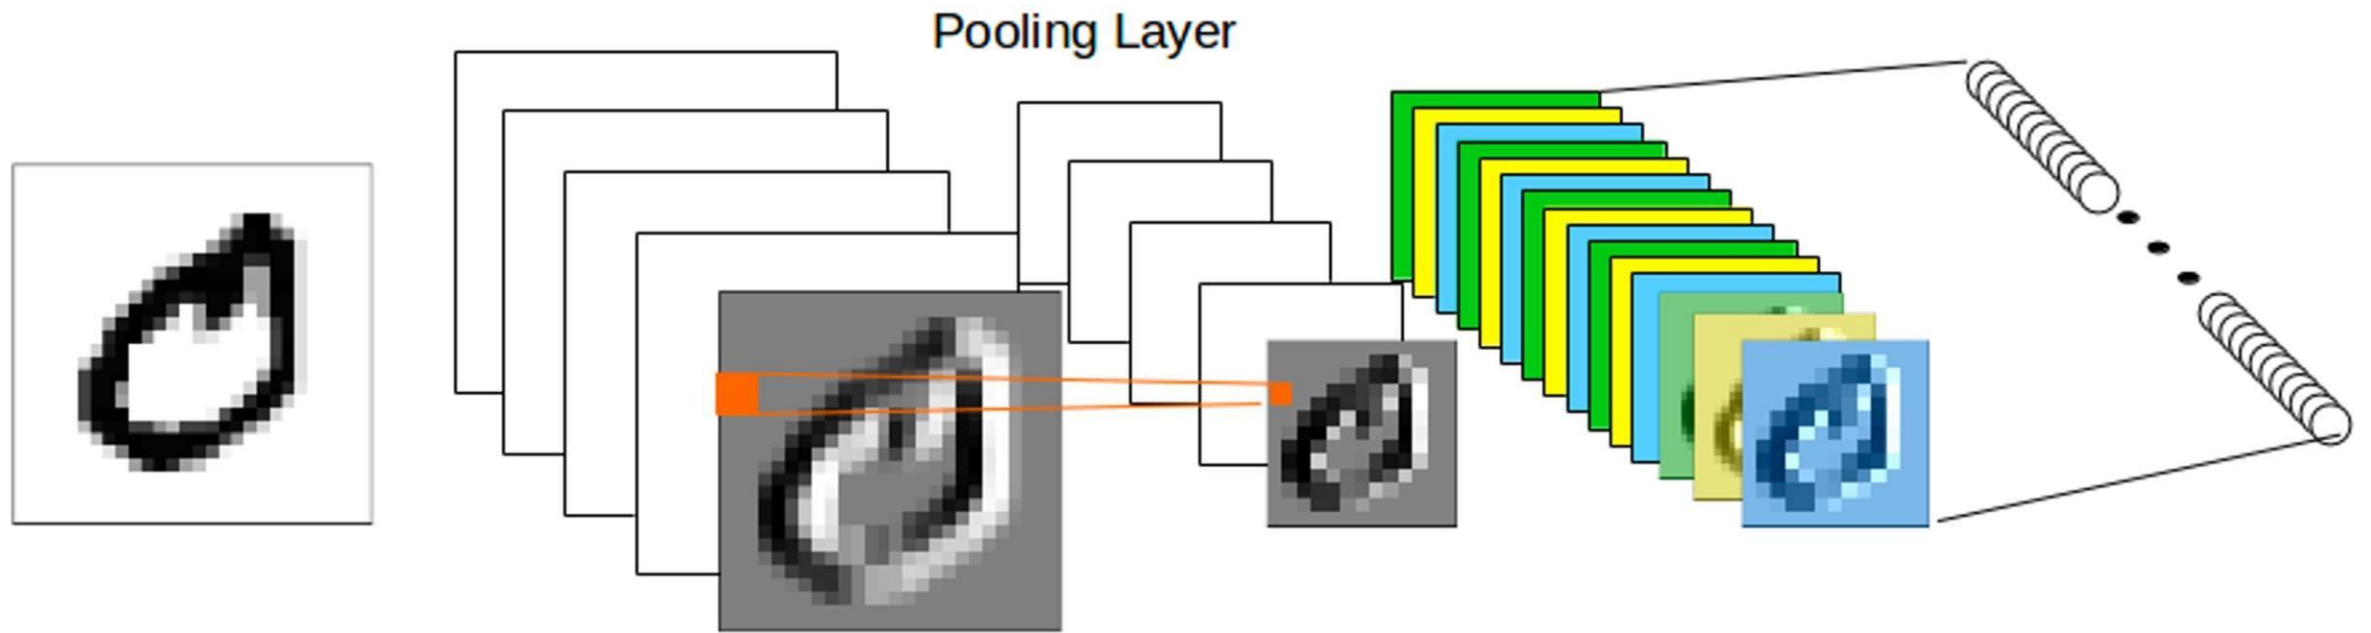
Convolutional Layer Parameters





Multiple Convolutional Layers



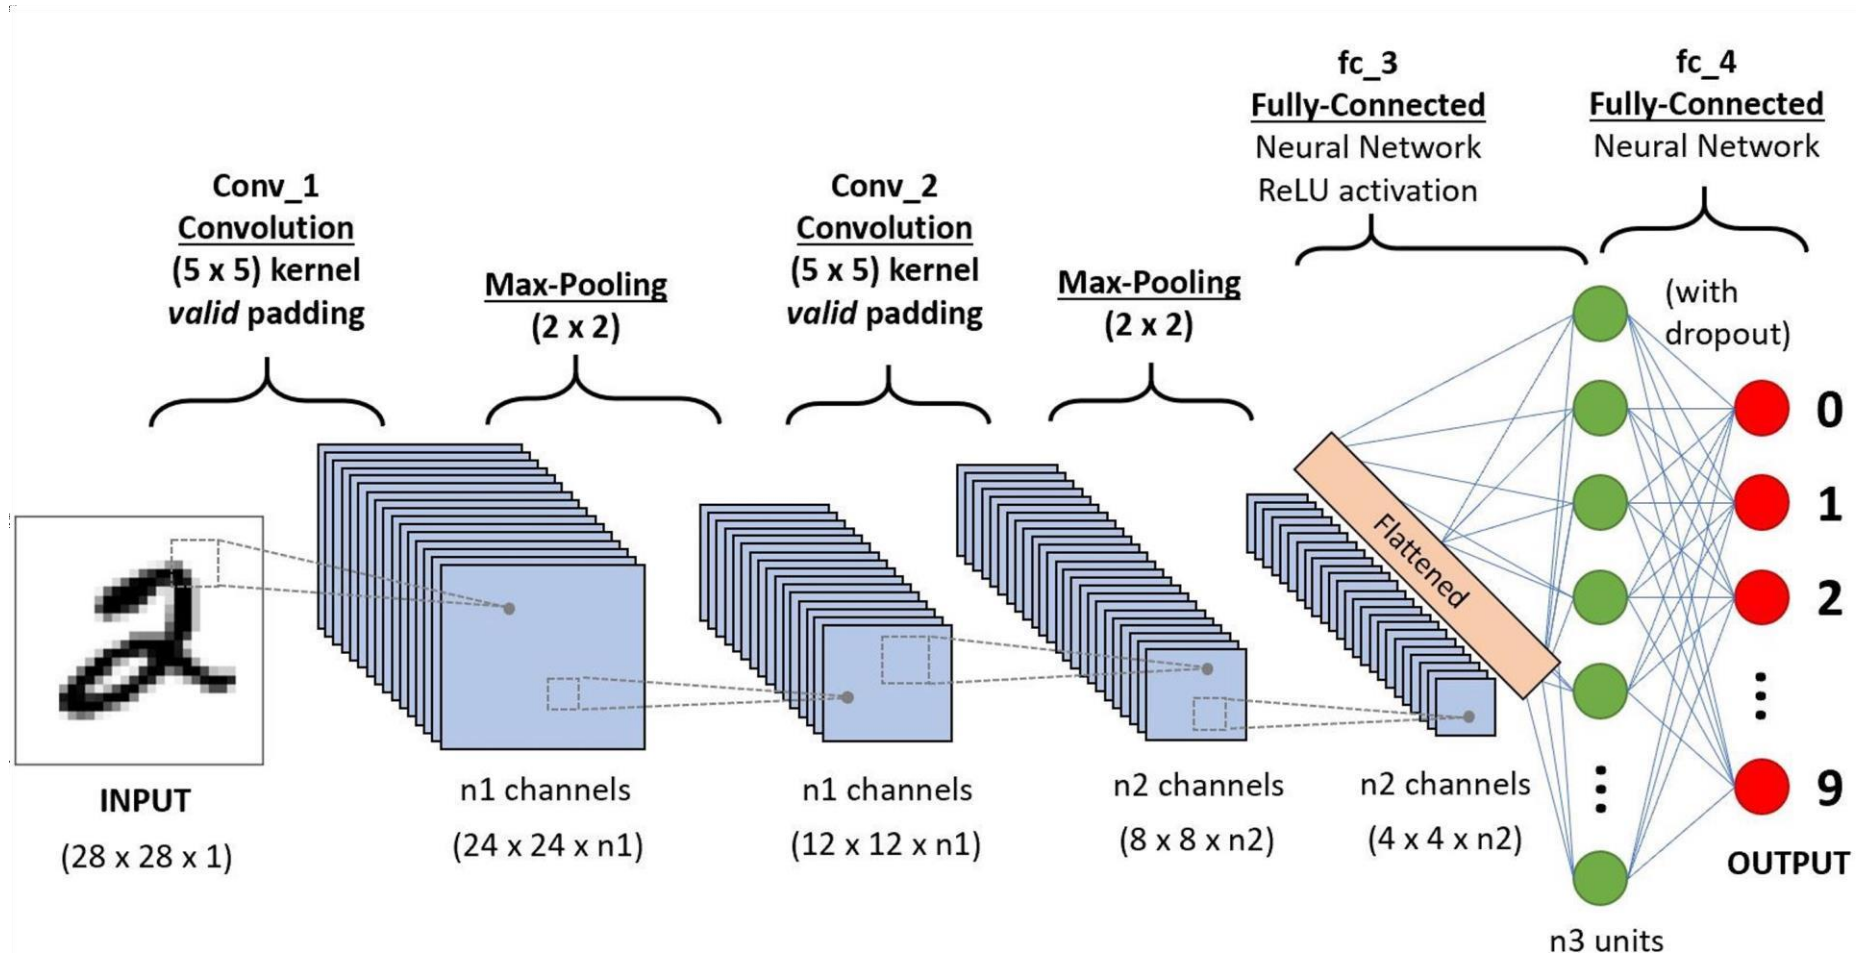
Pooling




Stride and Padding




CNNs for Image Processing



Basic Image Representation

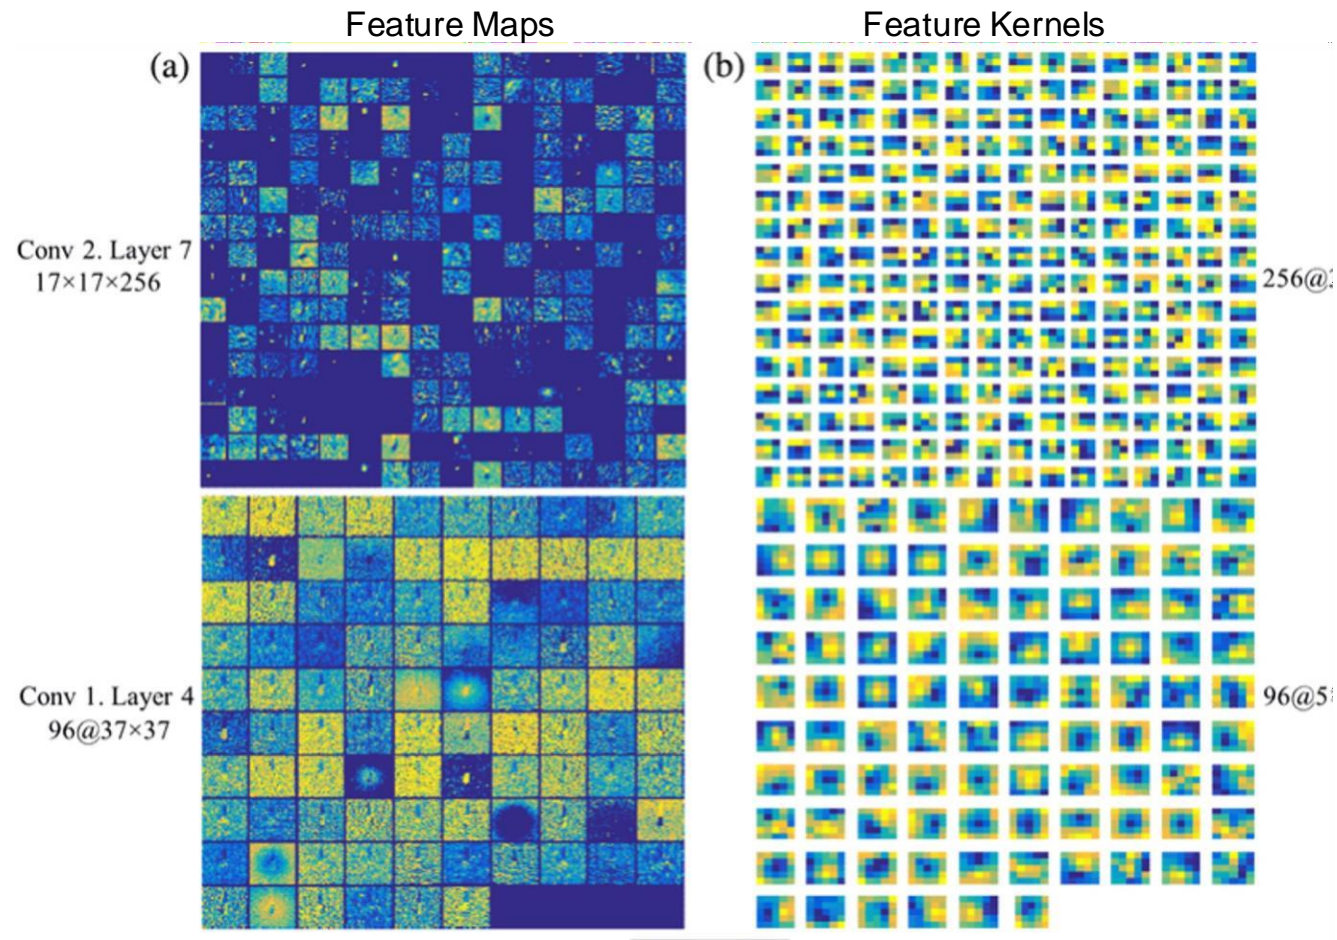


What do Feature Maps Look Like?




Learnings of First convolution layer on image of size 224X224X3


What do Feature Maps Look Like?




Comparison of different CNNs



Pre-Trained Weights



Sources of Training Data



Imagenet & MS COCO

Dataset examples

